

# Lecture one

## Introduction to MLIR and LLVM

# HPC motivation

- In HPC we rely heavily on libraries and compilers, and a number of these have already been ported on-top of MLIR, with more to follow
- But we need to be realistic here, MLIR was only created in 2022, so whilst it is growing very rapidly, it is still developing and being enhanced
  - This is one of the exciting aspects, as we in the HPC community have the opportunity to engage and influence its development
  - The ecosystem has matured significantly in the last 18 months, so it is now a realistic opportunity to build on-top of this and to expect benefits



TensorFlow



xDSL

# Technology overview

- MLIR is a framework for developing your own compiler Intermediate Representation (IR)
- Allowing different levels of abstraction to co-exist
  - Progressive lowering through these levels during compilation from one dialect to another
  - You the compiler/tool developer choose where these abstraction layers start & stop

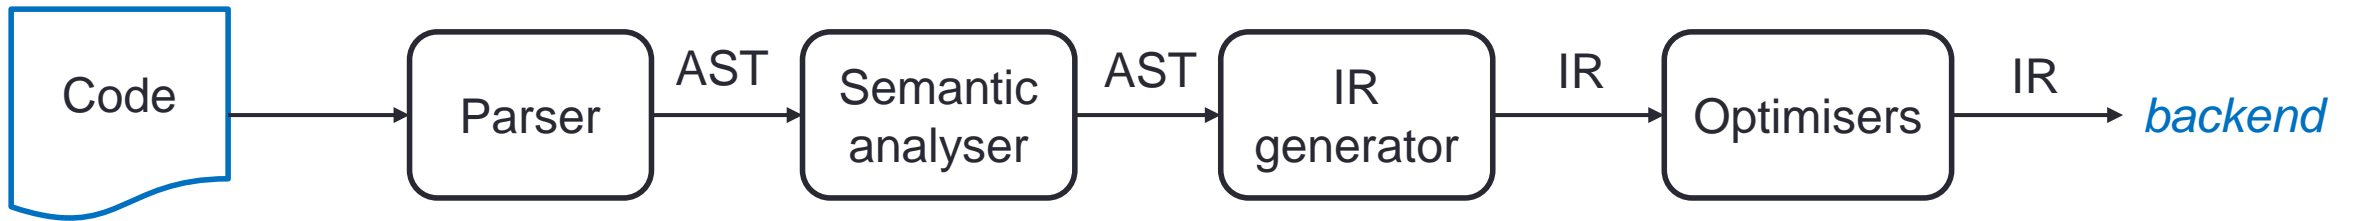


- LLVM is a collection of modular compiler technologies that aids reuse
  - Contains the clang, flang etc compilers
- For our uses, the most interesting part are the backends which target different architectures
- We can interact with these by generating LLVM-IR



# The challenge MLIR is looking to solve

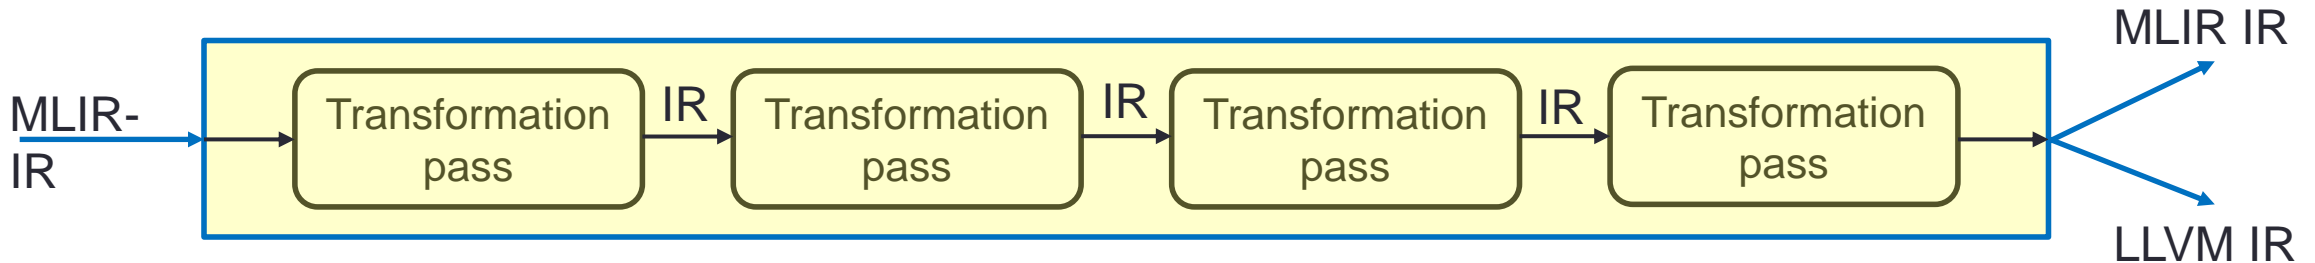
- Whilst LLVM has become very popular, much of the reuse is in the backend where LLVM-IR can target different architectures



- However, activities in the frontend can be extremely time consuming
  - Especially designing the IR, generating it and then and the optimisers
- This is where MLIR fits, providing a framework for defining and manipulating IRs
  - As well as many existing IRs and transformations to promote reuse

# Executing MLIR

- LLVM provides the *mlir-opt* tool which accepts inputs and generates outputs
  - Can input MLIR format IR, output can be MLIR IR, LLVM IR etc
  - Can instruct the tool to undertake passes on the IR



- Other tools wrap this, such as Flang which is LLVM's Fortran compiler
  - Parses the Fortran code into its own MLIR dialect (FIR) with other optional dialects, then undertakes transformation on the code before generating the object files

# MLIR-IR is human readable

- MLIR-IR can be viewed at each stage and is (fairly!) easy to read
  - Can be useful for debugging

```
func.func(%arg0) ({  
  %x = arith.constant() {"value" = 10 : i32} : () -> i32  
  %y = arith.addi(%x, %arg0) : (i32, i32) -> i32  
  func.return(%y) : (i32) -> i32  
}) {"sym_name" = "example"} (i32) -> (i32)
```

- This is Single Static Assignment (SSA) form, we will see this lots in the tutorial
- SSA comprises operations which have any number of operands (arguments), return zero, one or more results, and can contain static attributes
  - You can think of these a bit like function calls but translated by the compiler rather than each calling a function.

# MLIR-IR is the ecosystem

- MLIR provides the overarching framework here but known nothing about specific instructions

```
%x = "any_unknown_operation_here"(%arg0, %arg0) : (i32, i32) -> i32  
%y = "my_increment"(%x) : (i32) -> i32
```

- These will be carried through the transformation passes and it is expected that some of the transformations know how to handle these
- Therefore MLIR is an open, highly pluggable ecosystem where new dialects and new transformations can be created without having to modify others
- But there are standard dialects provided which include standard types

# Mixing dialects

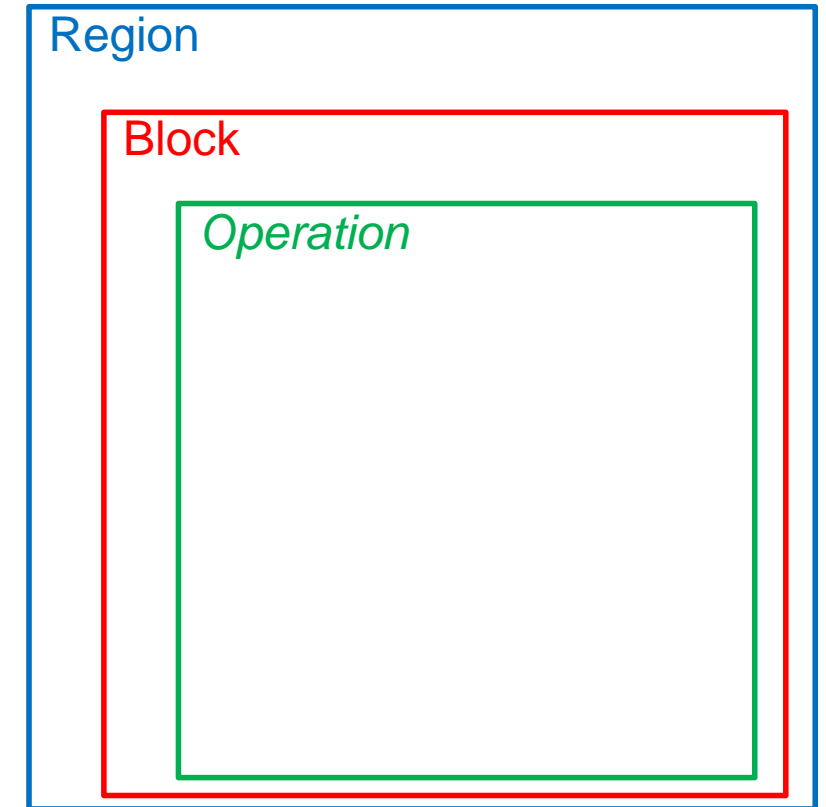
```
func.func(%arg0) ({  
  %x = arith.constant() {"value" = 10 : i32} : () -> i32  
  %y = arith.addi(%x, %arg0) : (i32, i32) -> i32  
  func.return(%y) : (i32) -> i32  
}) {"sym_name" = "example"} (i32) -> (i32)
```

- In this example we are already seeing some of MLIR's standard dialects being mixed and matched
  - The *func* dialect which marks this as a function and the return
  - The *arith* dialect that undertakes the integer addition (*addi*) and defines the constant (*constant*)
- Could add our own dialects and include them here too



# MLIR key concepts...

- MLIR is based on a graph style data structure of nodes which are called *Operations* and edges called *values*.
- Regions contain an ordered set of blocks, and blocks contain an ordered set of operations
  - Operations can also contain regions, and hence hierarchical structures can be expressed
- Operations are where most of the definition occurs
  - Can represent pretty much anything; functions definitions, function calls, variable allocation, conditionals, assignment etc
  - Operations can be arbitrarily extended



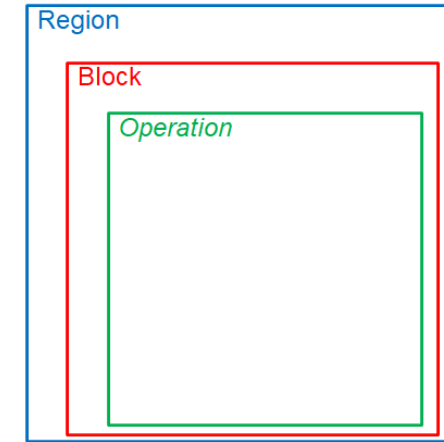
# Operations and their children

```
func.func(%arg0) ({  
  %x = arith.constant() {"value" = 10 : i32} : () -> i32  
  %y = arith.addi(%x, %arg0) : (i32, i32) -> i32  
  func.return(%y) : (i32) -> i32  
}) {"sym_name" = "example"} (i32) -> (i32)
```

- The *func* operation (from the *func* dialect) contains a region with a single block that itself contains a series of operations (the *constant*, *addi*, and *return* operations)
- The *func* operation also has a string attribute (*sym\_name*) which is set during compilation and an operand argument (*arg0*).

# Blocks

- Blocks contain a list of operations and can be thought of as executing these in order
- Blocks accept an optional list of arguments similar to a function



```
^bb0(%0: i32):  
  %1 = arith.constant() {"value" = 10 : i32} : () -> i32  
  %2 = arith.cmpi(%0, %1) {"predicate" = 2 : i32} : (i32, i32) -> i1  
  cf.cond_br(%2) [^bb1, ^bb2] { "operand_segment_sizes" = array<i32: 1, 0, 0> } : (i1)  
  
^bb1():  
  cf.br() [^bb2] : () -> ()  
  
^bb2():  
  .....
```

# Blocks

A block has zero, one or more arguments

Compare integers %0 and %1 for less than (2 passed to predicate) providing a 1 bit integer value (Boolean)

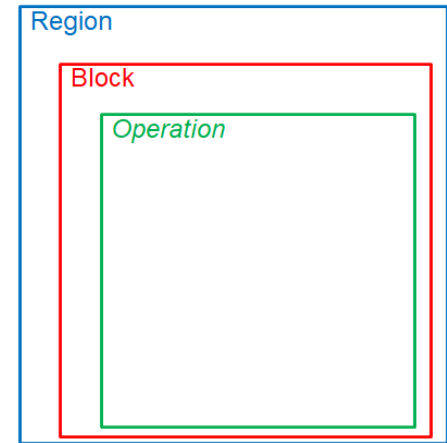
```
^bb0(%0: i32):  
  %1 = arith.constant() {"value" = 10 : i32} : () -> i32  
  %2 = arith.cmpi(%0, %1) {"predicate" = 2 : i32} : (i32, i32) -> i1  
  cf.cond_br(%2) [^bb1, ^bb2] { "operand_segment_sizes" = array<i32: 1, 0, 0> } : (i1)  
  
^bb1():  
  cf.br() [^bb2] : () -> ()  
  
^bb2():  
  .....
```

If true then jump to the `^bb1` block, otherwise jump to the `^bb2` block

Unconditional branch to block `^bb2`

# Regions

- A region contains an ordered list of blocks
- Regions must be contained within operations, therefore they are hierarchical
  - The body of a function is an example of a region, contained within the function definition operation
- Regions also provide encapsulation, where one can only branch to blocks within the same region



# Types

- MLIR provides its own type system and is strongly typed
  - MLIR provides the framework and the dialects provide the actual types themselves
    - Any number of types can be defined by dialects

```
func.func(%arg0) ({  
  %x = arith.constant() {"value" = 10 : i32} : () -> i32  
  %y = arith.addi(%x, %arg0) : (i32, i32) -> i32  
  func.return(%y) : (i32) -> i32  
}) {"sym_name" = "example"} (i32) -> (i32)
```

- Can also provide aliases
  - i.e. i32 is an alias for the IntegerType of width 32
  - f64 is an alias for the FloatType of width 64
  - Some types accept parameters – such as width and kind (useful for Fortran)

# Defining dialects

- Dialects are the mechanism by which the MLIR ecosystem is extended
- Can define new operations, new attributes, and new types
  - Each attribute has a unique name and this is prefixed to each of these
- Dialects are consumed by passes, with conversions between dialects
- Standard dialects provided as part of MLIR are:

- |                              |                        |                           |                              |
|------------------------------|------------------------|---------------------------|------------------------------|
| • <code>acc</code>           | • <code>complex</code> | • <code>ml_program</code> | • <code>sparse_tensor</code> |
| • <code>affine</code>        | • <code>dlti</code>    | • <code>nvgpu</code>      | • <code>tensor</code>        |
| • <code>amdgpu</code>        | • <code>emitc</code>   | • <code>nvvm</code>       | • <code>vector</code>        |
| • <code>amx</code>           | • <code>func</code>    | • <code>omp</code>        | • <code>x86vector</code>     |
| • <code>arith</code>         | • <code>gpu</code>     | • <code>pdl</code>        | • <code>Builtin</code>       |
| • <code>arm_neon</code>      | • <code>index</code>   | • <code>pdl_interp</code> | • <code>SPIR-V</code>        |
| • <code>arm_sve</code>       | • <code>linalg</code>  | • <code>quant</code>      | • <code>TOSA</code>          |
| • <code>async</code>         | • <code>LLVM</code>    | • <code>rocdl</code>      | • <code>Transform</code>     |
| • <code>bufferization</code> | • <code>math</code>    | • <code>scf</code>        |                              |
| • <code>cf</code>            | • <code>memref</code>  | • <code>shape</code>      |                              |

# Example flow.....

- We will see this in more detail in the hands-on activities, but to give you a feeling for it now

```
user@login01:~$ mlir-opt --convert-func-to-llvm ir.mlir | mlir-translate -mlir-to-llvmir | clang -x ir -o test -
```

Provide an MLIR file as input to the *mlir-opt* call which provides numerous transformations

Once in its final form the *mlir-translate* tool will convert from MLIR to a number of output formats, here we are generating LLVM IR

The standard LLVM IR is then provided to Clang which will compile it to an object file or executable

- But where do we get the MLIR input file (here *ir.mlir*) in the first place?
  - That's the job of our language specific frontend and we will see this in more detail in the hands on activities



# Conclusions

- MLIR is a very powerful framework, promoting reuse of infrastructure especially because it is backed by LLVM
  - You often see it presented as *doing properly* what from experience with LLVM they wish they had changed with that framework!
  - This is probably a little unfair to LLVM, but certainly it is based on all the lessons learnt
- We have focussed on the overarching concepts in this lecture
  - Regions, which contain blocks, which contain operations – which themselves can contain regions.
  - The ability to define dialects and transformations, and for these to coexist together
  - The fact that numerous standard dialects and transformations are provided *out-of-the-box*

