

Lecture two

The xDSL framework

MLIR is great, but....

- The MLIR framework has grown massively in popularity since it was introduced in 2020, with numerous projects relying upon it
 - Clearly it solves an important problem
- However, there are several disadvantages/challenges
 - Steep learning curve due to it being written in C++ and Cmake
 - Requirement to use esoteric Tablegen DSL to define dialects
 - Need to download and build the entirety of LLVM (it's big!)
 - API is fairly complex and not particularly well documented
 - A moving target, with lots of changes happening



Step in xDSL.....

xDSL

- xDSL provides the ability to work with MLIR in Python
 - But crucially is far more than simply a Python wrapper!
- Contains the full MLIR functionality and concepts but with significant enhancements aimed at easing development of dialects and transformations
 - IRDL is our way of defining dialects, with tools to generate to/from Tablegen
 - Therefore, dialects are also written in Python and much more descriptive than in MLIR
 - Using a high productivity programming language one can easily experiment
 - No need to install LLVM, just need to grab the *xds/* Python package from pip
 - Provides the standard MLIR dialects (and many others!)
 - Compatibility with MLIR, where output of MLIR can be consumed by xDSL and vica-versa

A burgeoning community

- BSD licenced and available at <https://github.com/xdslproject/xdsl>
- Alternatively, can install the latest release version via pip
 - Releases happen every few weeks
 - Currently version 0.11
- Lots of activity on the Github repository, and a chatty community on Zulip <https://xdsl.zulipchat.com/>

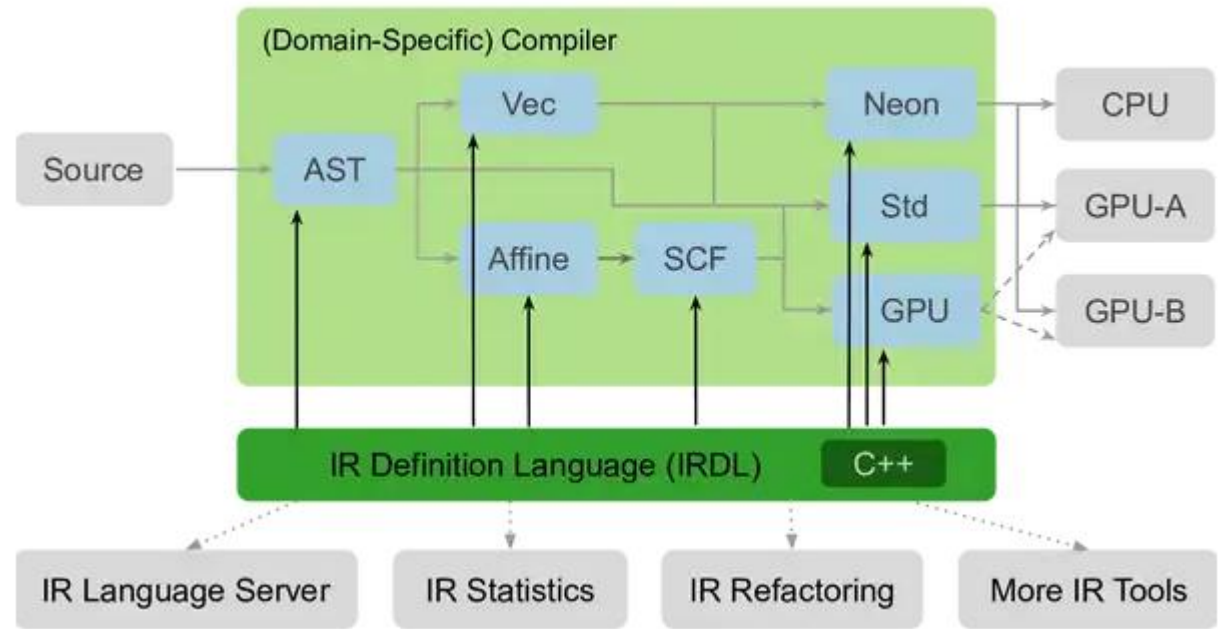
- Therefore, at the end of today's session there is a community you can connect with and we will work with you help integrate xDSL with your code
 - We will share all the links at the end of the session

The screenshot displays the GitHub repository for 'georgemitenkov frontend: Add basic for loop support (#800)'. The repository has 795 commits and is a Python Compiler Design Toolkit. The commit history table lists recent changes, including CI updates, pre-commit hooks, and documentation. The right sidebar shows repository statistics: 62 stars, 14 watching, and 27 forks. It also lists 18 releases, with the latest being v0.11.0 from last week, and 2 packages: xdsl/mlir-cache and xdsl/mlir. The 'Used by' section shows 10 users, and the 'Contributors' section shows 27 contributors.

File	Commit Message	Time Ago
.github	CI: Change CI name for better readability (#786)	2 days ago
bench/parser	CI: switch formatter to black. (#763)	5 days ago
docs	CI: Add pre-commit for black, whitespaces, and end-of-line (#775)	2 days ago
tests	frontend: Add basic for loop support (#800)	39 minutes ago
xdsl	frontend: Add basic for loop support (#800)	39 minutes ago
.coveragerc	misc: add Toy chapter 1 python code, examples and notebook (#354)	3 months ago
.git-blame-ignored-commits	CI: switch formatter to black. (#763)	5 days ago
.gitignore	docs: Init xDSL-MLIR interoperation tutorial	last month
.pre-commit-config.yaml	CI: Add pre-commit for black, whitespaces, and end-of-line (#775)	2 days ago
CHANGELOG.md	CI: Add pre-commit for black, whitespaces, and end-of-line (#775)	2 days ago
LICENSE	Create python package	2 years ago
MANIFEST.in	install: Add typing stubs in the PyPi install (#223)	5 months ago
README.md	CI: Add pre-commit for black, whitespaces, and end-of-line (#775)	2 days ago
codecov.yml	CI: Upload the line by line coverage (#417)	2 months ago
pyproject.toml	misc: lint test_xdsl_opt.py (#789)	2 days ago
requirements-optional.txt	pip prod(deps): bump pyright from 1.1.304 to 1.1.305 (#792)	yesterday
requirements.txt	CI: Add pre-commit for black, whitespaces, and end-of-line (#775)	2 days ago
setup.cfg	CI: Add pre-commit for black, whitespaces, and end-of-line (#775)	2 days ago
setup.py	CI: switch formatter to black. (#763)	5 days ago
versioneer.py	CI: switch formatter to black. (#763)	5 days ago

A key component: IRDL

- In MLIR one has to use the esoteric Tablegen to define dialects
- Instead, we provide IRDL which is a DSL for describing dialects using the same concepts (regions, blocks, operations) as MLIR but in a more straightforward manner
- Crucially, IRDL is fully compatible with MLIR via a C++ implementation, so the definition of dialects does not diverge from MLIR itself.
- We will not go into IRDL in detail, instead informally presenting this to you by looking at the Python code



Defining a dialect

- Let's consider the IR node representing a function definition – this is the definition of the node from the hands-on exercise that we will look at soon

```
@irdl_op_definition
class Function(IRDLOperation):
    """
    A Python function, our handling here is simplistic and limited but sufficient
    for the exercise (and keeps this simple!) You can see how we have a mixture of
    attributes and a region for the body
    """
    name = "tiny_py.function"

    fn_name: OpAttr[StringAttr]
    args: OpAttr[ArrayAttr]
    return_var: OpAttr[AnyAttr()]
    body: Region

    @staticmethod
    def get(fn_name: str | StringAttr,
            return_var: Operation | None,
            args: List[Operation],
            body: List[Operation],
            verify_op: bool = True) -> Routine:
        if isinstance(fn_name, str):
            # If fn_name is a string then wrap it in StringAttr
            fn_name = StringAttr(fn_name)

        if return_var is None:
            # If return is None then use the empty token placeholder
            return_var = EmptyType()
        res = Function.build(attributes={"fn_name": fn_name, "return_var": return_var,
                                       "args": ArrayAttr(args)}, regions=[Region([Block(body)])])
        if verify_op:
            # We don't verify nested operations since they might have already been verified
            res.verify(verify_nested_ops=False)
        return res
```

Decorator and inheritance defines the class to be an IRDL operator

These are attributes associated with the IR node, you can see that the function name is a String, the arguments are an array and the return var is any type

Constructs the Function operation based on the parameters provided. Can also call *build* directly or the *create* which is similar

Inbuilt verification ensuring operation correctly created

The name identifies this in the IR output

This is a region (remember operations can nest regions)

Generating the IR

```
def visit_FunctionDef(self, node):
    contents=[]
    for a in node.body:
        contents.append(self.visit(a))
    return tiny_py.Function.get(node.name, None, [], contents)
```

- Our, extremely simple, parser will create the IR node like so, where we provide it with an array representing the contents, and empty array for arguments and the name

```
builtin.module() {
  tiny_py.module() {
    tiny_py.function() ["fn_name" = "hello_world", "return_var" = !tiny_py.emptytoken, "args" = []] {
      ....
    }
  }
}
```

- Just as we have seen with MLIR in the previous lecture, the output generated from this is human readable
 - Can see we are using the EmptyToken to represent no return variable, this is defined within the dialect itself and the semantics of this is constrained within the dialects and aware transformations

The xDSL opt tool

- In the hands on exercise we will see the *tinypy-opt* tool which drives the transformations on IRs
 - These opt tools are standard in MLIR/LLVM, and xDSL provides Python infrastructure to make it trivial to write these for specific uses

```
user@login01:~$ ./tinypy-opt code.xdsl -t mlir
```

- We are instructing xDSL to parse *code.xdsl* contents & output it in MLIR format
 - This is different to the output on the previous slide, as that is in xDSL format (which we think is clearer). However, MLIR format is required to interoperate with MLIR.

```
"builtin.module"() ({
  "tiny_py.module"() ({
    "tiny_py.function"() ({
      ....
    }) {"fn_name" = "hello_world", "return_var" = #tiny_py.emptytoken, "args" = []} : () -> ()
  }) : () -> ()
}) : () -> ()
```


Undertaking transformations on the IR

- Transformations are written in Python and driven by the *opt* tool

```
user@login01:~$ ./tinypy-opt code.xdsl -p tiny-py-to-standard
```

- This transformation converts the tiny-py dialect to standard dialects
 - Here converting from `tiny_py.function` into `func.func`
 - This is an example of lowering from a dialect with a richer source of information to lower level standard dialects

```
builtin.module() {  
  func.func() ["sym_name" = "hello_world", "function_type" = !fun<[], []>, "sym_visibility" = "public"] {  
    ....  
    func.return()  
  }  
}
```

Lifting the lid on transformations

Entry point of the transformation

```
@dataclass
class LowerTinyPyToStandard(ModulePass):

    name = 'tiny-py-to-standard'

    def apply(self, ctx: MLContext, input_module: ModuleOp):
        res_module = translate_program(input_module)
        res_module.regions[0].move_blocks(input_module.regions[0])
```

Pass in the SSA context which maps all variable names to their SSA value

Routine that translates the tiny_py assign node to the standard dialects

```
def translate_assign(ctx: SSAValueCtx,
                    assign: tiny_py.Assign) -> List[Operation]:

    """
    Translates assignment
    """
    var_name = assign.var_name
    assert isinstance(var_name, StringAttr)

    expr, ssa=translate_expr(ctx, assign.value.blocks[0].ops[0])

    # Always update the SSA context as it is this new SSA element that subsequent references
    # to the variable should reference
    ctx[var_name] = ssa
    return expr
```

Extract the name from the tiny py dialect and check this is a string attribute

Translate the RHS (which is an expression) and this returns both the operations and SSA value

In the SSA context map this variable name to the expression SSA value

Manipulating the IR

- We also provide several ways of manipulating and rewriting the IR, these are based on MLIR standard approaches (but much easier in Python!)
 - This is registered with the opt tool and executed via `-p apply-my-analysis`

Obtain the operation's parent block

Detach this node from the IR

Insert this new operation in the block at index *idx*

```
class ApplyRewriter(RewritePattern):
    @op_type_rewrite_pattern
    def match_and_rewrite(
        self, call_node: tiny_py.CallExpr, rewriter: PatternRewriter):
        block = call_node.parent
        idx = block.ops.index(call_node)
        call_node.detach()
        ...
        some_other_op = ....
        rewriter.insert_op_at_pos(some_other_op, block, idx)

def apply_my_analysis(ctx: psy_ir.MLContext, module: ModuleOp) -> ModuleOp:
    applyRewriter=ApplyRewriter()
    walker = PatternRewriteWalker(GreedyRewritePatternApplier([applyRewriter]), apply_recursively=False)
    walker.rewrite_module(module)

    return module
```

Will match (and execute this function) for all nodes in the IR that are of type *tiny_py.CallExpr*

Get the index of this node in the block's list of operations

Create some other operation, e.g. containing the *call_node* operation

Entry point of the pass, we can apply many transformations here if we wish

- All these functions, and many more, to manipulate the IR are documented

Different levels of IR

```
builtin.module() {  
  tiny_py.module() {  
    tiny_py.function() ["fn_name" = "ex_two",  
      "return_var" = !empty, "args" = []] {  
      tiny_py.assign() ["var_name" = "val"] {  
        tiny_py.constant() ["value" = 0.0 : !f32]  
      }  
      tiny_py.assign() ["var_name" = "add_val"] {  
        tiny_py.constant() ["value" = 88.2 : !f32]  
      }  
      tiny_py.loop() ["variable" = "a"] {  
        tiny_py.constant() ["value" = 0 : !i32]  
      } {  
        tiny_py.constant() ["value" = 100000 : !i32]  
      } {  
        tiny_py.assign() ["var_name" = "val"] {  
          tiny_py.binaryoperation() ["op" = "add"] {  
            tiny_py.var() ["variable" = "val"]  
          } {  
            tiny_py.var() ["variable" = "add_val"]  
          }  
        }  
      }  
    }  
  }  
}
```

- Hierarchical
- Close representation to the original code
- Easy to reason about in transformation passes



*Can transform
between these,
also possible to
mix levels*

```
"builtin.module"() ({  
  "func.func"() ({  
    %0 = "arith.constant"() {"value" = 0.0 : f32} : () -> f32  
    %1 = "arith.constant"() {"value" = 88.2 : f32} : () -> f32  
    %2 = "arith.constant"() {"value" = 0 : i32} : () -> i32  
    %3 = "scf.while"(%2, %0) ({  
      %4 = "arith.constant"() {"value" = 100000 : i32} : () -> i32  
      %5 = "arith.cmpi"(%4, %2) {"predicate" = 1 : i64} : (i32, i32) -> i1  
      "scf.condition"(%5, %2) : (i1, i32) -> ()  
    }, {  
      ^0(%6 : i32, %7 : f32):  
        %8 = "arith.addf"(%0, %1) : (f32, f32) -> f32  
        %9 = "arith.constant"() {"value" = 1 : i32} : () -> i32  
        %10 = "arith.addi"(%2, %9) : (i32, i32) -> i32  
        "scf.yield"(%10, %8) : (i32, f32) -> ()  
      }) : (i32, f32) -> i32  
    "func.return"() : () -> ()  
  }) {"sym_name" = "ex_two", "function_type" = () -> (), "sym_visibility"  
= "public"} : () -> ()  
}) : () -> ()
```

- SSA form, much closer to concrete implementation
- Exposes lower level details and concerns
- More difficult to see how relates to original code

Which level to run transformations on?

- It really depends on the transformation!
- Some will suit working on the higher level IR much more, for instance discovering parallelism or mapping code constructs to features such as stencils

```
"builtin.module"() ({  
  "func.func"() ({  
    %0 = "arith.constant"() {"value" = 0.0 : f32} : () -> f32  
    %1 = "arith.constant"() {"value" = 88.2 : f32} : () -> f32  
    %2 = "arith.constant"() {"value" = 0 : i32} : () -> i32  
    %3 = "scf.while"(%2, %0) ({  
      %4 = "arith.constant"() {"value" = 100000 : i32} : () -> i32  
      %5 = "arith.cmpi"(%4, %2) {"predicate" = 1 : i64} : (i32, i32) -> i1  
      "scf.condition"(%5, %2) : (i1, i32) -> ()  
    }, {  
      ^0(%6 : i32, %7 : f32):  
        %8 = "arith.addf"(%0, %1) : (f32, f32) -> f32  
        %9 = "arith.constant"() {"value" = 1 : i32} : () -> i32  
        %10 = "arith.addi"(%2, %9) : (i32, i32) -> i32  
        "scf.yield"(%10, %8) : (i32, f32) -> ()  
    }) : (i32, f32) -> i32  
    "func.return"() : () -> ()  
  }) {"sym_name" = "ex_two", "function_type" = () -> (), "sym_visibility"  
= "public"} : () -> ()  
}) : () -> ()
```

```
builtin.module() {  
  tiny_py.module() {  
    tiny_py.function() {"fn_name" = "ex_two",  
      "return_var" = !empty, "args" = []} {  
      tiny_py.assign() {"var_name" = "val"} {  
        tiny_py.constant() {"value" = 0.0 : !f32}  
      }  
      tiny_py.assign() {"var_name" = "add_val"} {  
        tiny_py.constant() {"value" = 88.2 : !f32}  
      }  
      tiny_py.loop() {"variable" = "a"} {  
        tiny_py.constant() {"value" = 0 : !i32}  
      } {  
        tiny_py.constant() {"value" = 100000 : !i32}  
      } {  
        tiny_py.assign() {"var_name" = "val"} {  
          tiny_py.binaryoperation() {"op" = "add"} {  
            tiny_py.var() {"variable" = "val"}  
          } {  
            tiny_py.var() {"variable" = "add_val"}  
          }  
        }  
      }  
    }  
  }  
}
```

- Others suit working on the lower level IR, such as constant folding
- The point is xDSL/MLIR is very powerful and gives you the choice
 - Select what works best for your dialects and code!

Documentation

- In this lecture we have provided a general overview of how the different parts of xDSL fit together and can be used as part of the compilation flow
 - We will see these in the hands-on exercises and gain more experience with the different aspects
- However there is much more depth than we would want to cover in an introduction tutorial, at <https://github.com/xdslproject/xdsl/tree/main/docs> we have notebooks that you can explore afterwards to delve deeper
 - *xdsl-introduction.ipynb* provides an overview of xDSL and key concepts
 - *database_example.ipynb* walks you through the creation of a database DSL using xDSL alone (e.g. it never calls into MLIR or LLVM) that generates SQL for querying databases
 - The *toy* directory is the MLIR toy example in xDSL

Conclusions



- xDSL enables us to leverage MLIR constructs in a high productivity Python environment
- Don't need MLIR or LLVM installed to use xDSL
 - Although will need it if you want to generate LLVM-IR and use the LLVM backends
- Lots of resources on xDSL that you can refer to
 - Website: <https://xdsl.dev/>
 - Zulip chat: <https://xdsl.zulipchat.com/>
 - Jupyter notebooks: <https://github.com/xdslproject/xdsl/tree/main/docs>
- A number of users of xDSL in the community
 - Devito
 - PSyclone
 - CONVOLVE

